

# Speech Analysis and Feature Extraction

TTT4185 Machine Learning for Signal Processing

Giampiero Salvi

Department of Electronic Systems  
NTNU

HT2020


## 1 Speech Signal Representations

- Signal Processing Reminder
- Sampling and Quantization
- Pre-Emphasis
- Windowing
- Discrete Fourier Transform

## 2 Feature Extraction


- Linear Prediction Analysis (LPA)
- Cepstrum
- Perceptually Motivated Features

# Assuming that you know



- linear time-invariant systems (continuous and discrete time case)
- convolution
- impulse response
- Fourier transform and transfer function

# Convolution and Impulse Response



$$y[n] = T\{x[n]\} = \sum_{k=-\infty}^{\infty} x[k]h[n-k] = x[n] * h[n]$$


Properties:

- linearity:  $(a_1x_1 + a_2x_2) * h = a_1(x_1 * h) + a_2(x_2 * h)$
- symmetry:  $x * h = h * x$

Kind of complicated to interpret.



# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



# Convolution: Illustration



$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$







# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



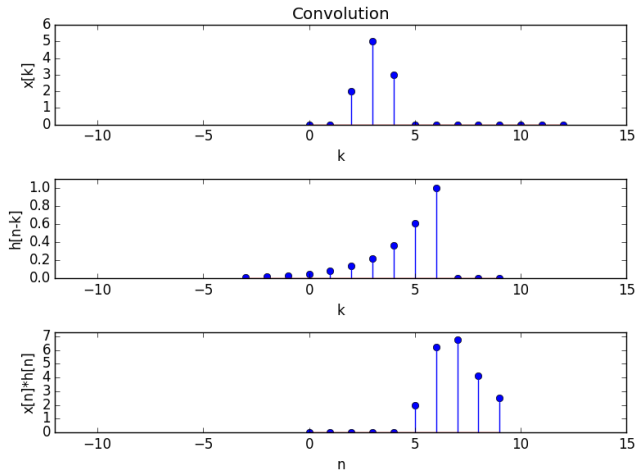

# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



# Convolution: Illustration

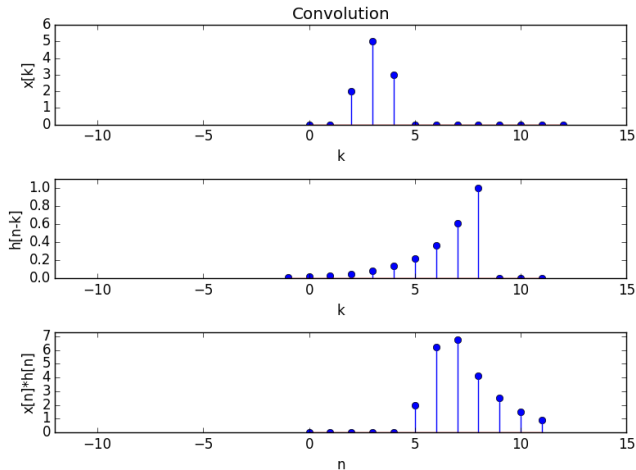

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





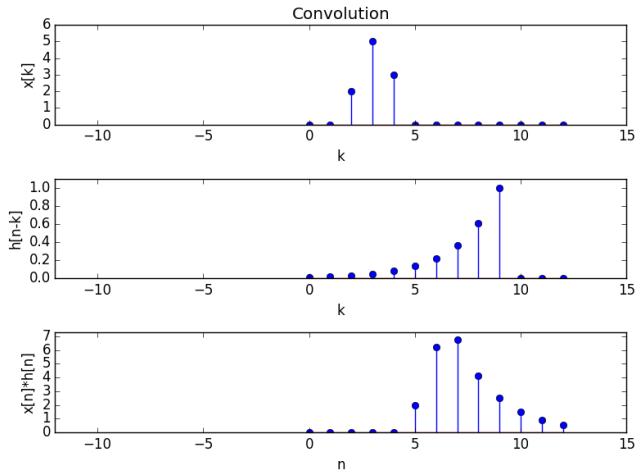

# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



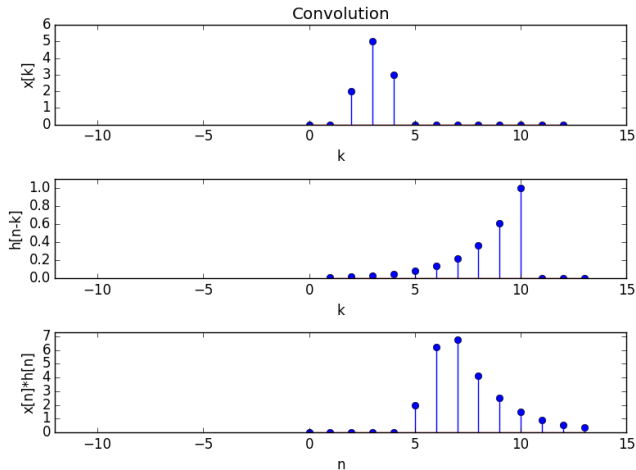

# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



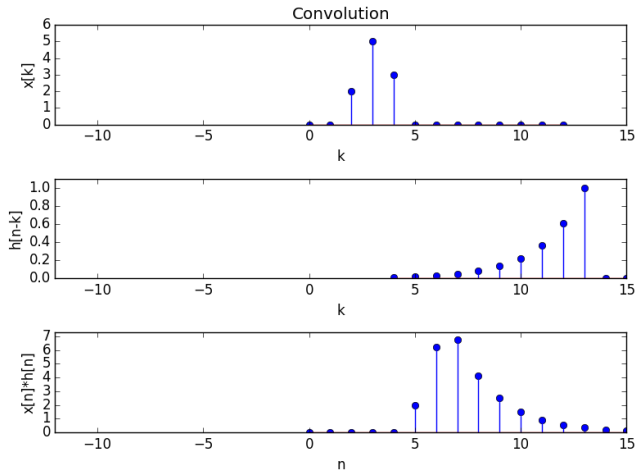

# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



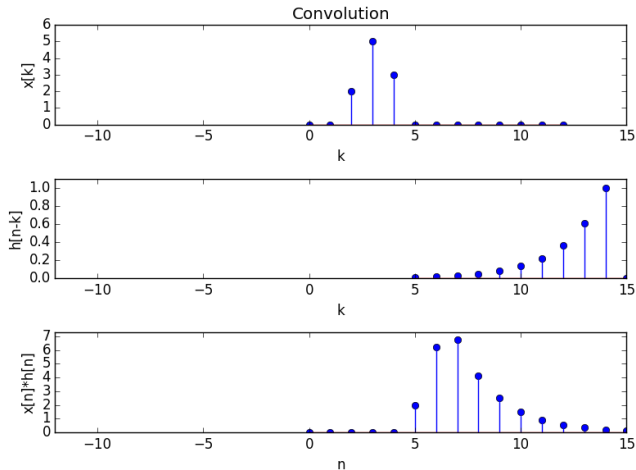

# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



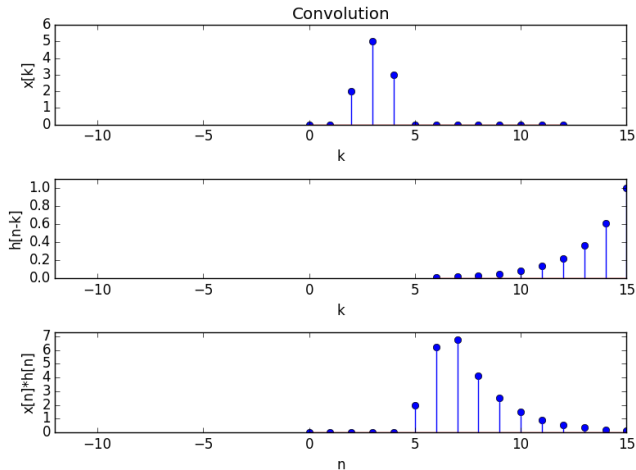

# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$



# Convolution: Illustration

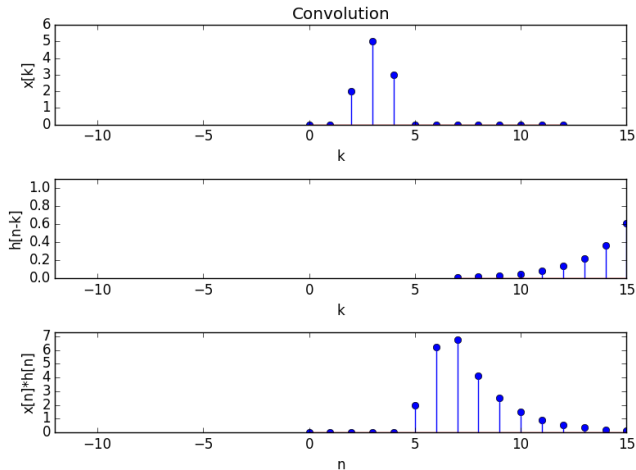

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$





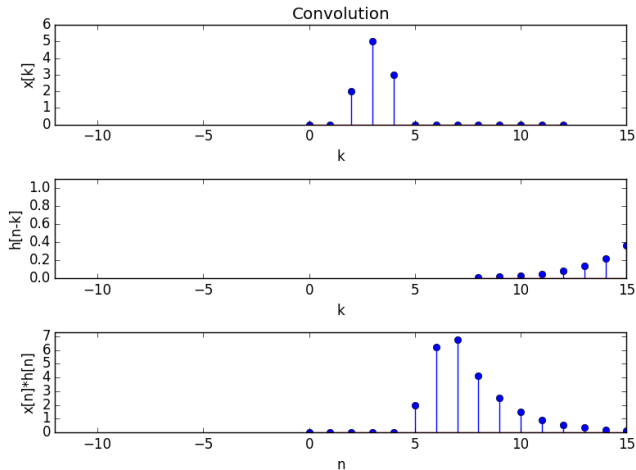

# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$




# Convolution: Illustration

$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$




# Question: Which is the convolution of the signals


$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$




a)




b)



c)




d)




# Question: Which is the convolution of the signals


$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$




a)




b)



c)




d)




# Question: Which is the convolution of the signals


$$x[n] * h[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k]$$




a)




b)



c)



d)



# Linear filter in general (Infinite Impulse Response)

`scipy.signal.lfilter(b, a, x, ...)`

$$\begin{aligned}y[n] &= \frac{1}{a_0} \left( b_0 x[n] + b_1 x[n-1] + \dots + b_P x[n-P] + \right. \\ &\quad \left. - a_1 y[n-1] - a_2 y[n-2] - \dots + a_Q y[n-Q] \right) \\ &= \frac{1}{a_0} \left( \sum_{i=0}^P b_i x[n-i] - \sum_{j=1}^Q a_j y[n-j] \right)\end{aligned}$$

$$a = [a_0, a_1, \dots, a_Q]$$

$$b = [b_0, b_1, \dots, b_P]$$

# Fourier Transforms

Fourier transform of continuous signals

$$X(\omega) = \int_{-\infty}^{\infty} x(t)e^{-j\omega t} dt$$

Fourier transform of discrete signals

$$X(\omega) = \sum_{k=-\infty}^{\infty} x[k]e^{-j\omega k}$$

Discrete Fourier Transform (and Fast Fourier Transform)

$$X[n] = \sum_{k=0}^{N-1} x[k]e^{-j2\pi \frac{n}{N}k}$$


# Properties of Fourier Transform

Property	Time domain		Frequency domain
Linearity	$ax[n] + by[n]$	$\iff$	$aX(\omega) + bY(\omega)$
Time-shift	$x[n - k]$	$\iff$	$X(\omega)e^{-j\omega nT_s}$
Convolution	$x[n] * h[n]$	$\iff$	$X(\omega)H(\omega)$
Multiplication	$x[n]w[n]$	$\iff$	$X(\omega) * W(\omega)$




# Fourier Transform of Useful Signals

Sinusoidal at frequency  $f = 5\text{Hz}$




Single (Ideal) Impulse




# Fourier Transform of Useful Signals

## Train of (Ideal) Impulses



## Square function



# Properties of Fourier Transform — Example: sampling


Time domain

Frequency domain


$$x[n]w[n] \iff X(\omega) * W(\omega)$$

TIME DOMAIN


FREQUENCY DOMAIN



# Properties of Fourier Transform — Example: sampling



# Sampling Theorem (Nyquist-Shannon)



If  $x(t)$  contains energy up to  $B_x$ , in order to reconstruct the signal we need to sample with

$$f_s > 2B_x$$

# Aliasing




Figure from Huang, Acero and Hon (2001)

# Aliasing: Illustration



Video from <https://youtu.be/usN47Jvy9PY> (unfortunately removed)

# First step: represent speech signal

## Sampling

- **Nyquist-Shannon Theorem:** sample at twice the band
- 8kHz (4kHz band, telephone), 16kHz (8 kHz band, high quality)
- TIDIGITS sampled at 20kHz
- TIMIT sampled at 16kHz

## Quantisation


- Type of quantisation: linear, a-law,  $\mu$ -law
- 8, 16 bits (more rare 32, floating point)
- TIDIGITS and TIMIT are quantised with 16 bits linear



# Pre-emphasis

Compensate for the 6db/octave drop (glottal shape - radiation at the lips)

$$y[n] = x[n] - \alpha x[n - 1]$$



$\alpha$  is usually 0.95–0.97



# Pre-Emphasis as Linear Filter

Time domain


Frequency domain

$$x[n] * h[n] \iff X(\omega)H(\omega)$$

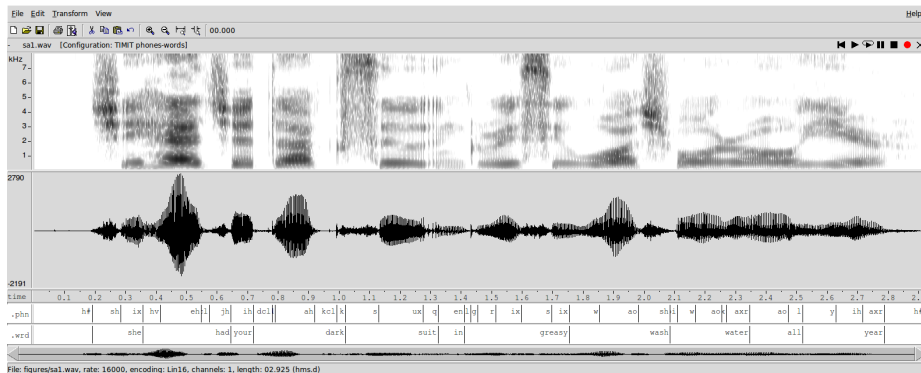
$$y[n] = x[n] - \alpha x[n-1], \quad \text{with } \alpha = 0.97$$



# Pre-emphasis applied to vowel




# A time varying signal





- speech is time varying
- short segments are quasi-stationary
- use short time analysis



# Short-Time Fourier Analysis





# Short-Time Fourier Analysis



# Short-Time Fourier Analysis



# Short-Time Fourier Analysis





# Properties of Fourier Transform — Example: Windowing


Time domain

Frequency domain

$$x[n]w[n]$$


$$\Longleftrightarrow$$

$$X(\omega) * W(\omega)$$




# Effect of Windowing

## Square window




## Hamming window





# Effect of Windowing on Speech

## Effect of different window functions





Rectangular  
30 ms



Rectangular  
15 ms

Hamming  
30 ms



Hamming  
15 ms


# Windowing, typical values

- signal sampling frequency: 8–20kHz
- analysis window: 10–50ms
- frame step: 10–25ms (100–40Hz)

# Fast Fourier Transform (FFT) ( $N = 512$ )

`scipy.fftpack.fft(x, n=512, ...)`


$$X[n] = \sum_{k=0}^{N-1} x[k] e^{-j2\pi \frac{n}{N} k}$$



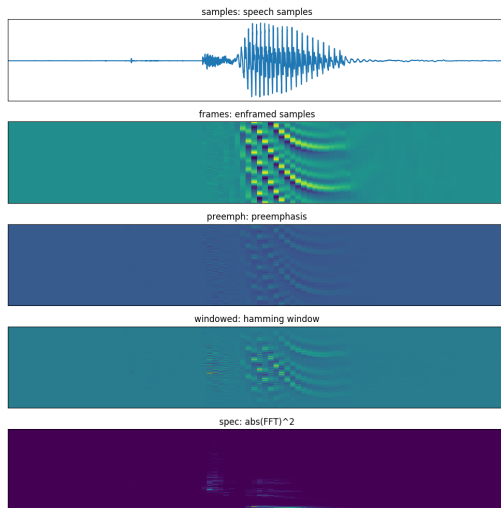
# Fast Fourier Transform (FFT) ( $N = 512$ )

`scipy.fftpack.fft(x, n=512, ...)`

$$X[n] = \sum_{k=0}^{N-1} x[k] e^{-j2\pi \frac{n}{N} k}$$



# Speech signal processing in practice



## 1 Speech Signal Representations

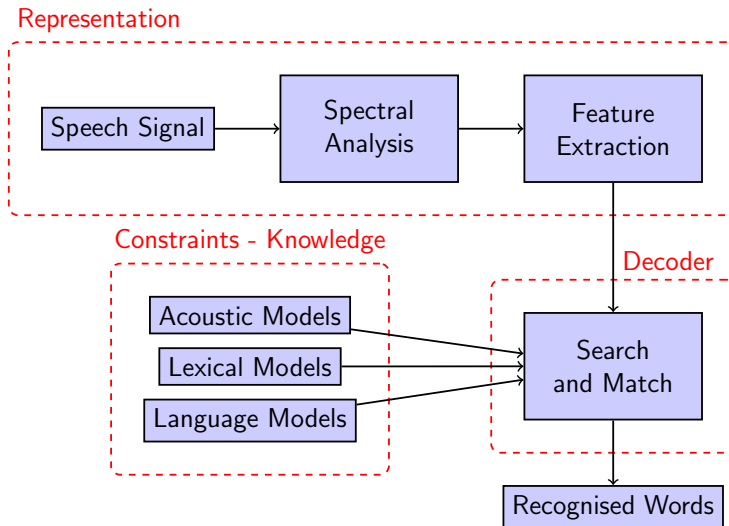
- Signal Processing Reminder
- Sampling and Quantization
- Pre-Emphasis
- Windowing
- Discrete Fourier Transform

## 2 Feature Extraction

- Linear Prediction Analysis (LPA)
- Cepstrum
- Perceptually Motivated Features




# Components of ASR System



# Speech Signal Representations

## Representation




## Goals:

- disregard irrelevant information
- optimise relevant information for modelling

# Speech Signal Representations

## Representation




Means:

- try to model essential aspects of speech production
- imitate auditory processes
- consider properties of statistical modelling


# $F_0$ and Formants

- Varying  $F_0$  (vocal fold oscillation rate)

spectrum (log)  $f_0 = 100\text{Hz}$




spectrum (log)  $f_0 = 250\text{Hz}$




- Varying Formants (vocal tract shape)


spectrum (log) vowel [ɛ]



spectrum (log) vowel [u]



# Linear Prediction Coefficients (LPC)



approximate  $y[n]$  as a linear combination of  $p$  previous samples:

$$\hat{y}[n] = \sum_{k=1}^p a_k y[n - k]$$

The error is called **residual**:  $r[n] = \hat{y}[n] - y[n]$


The output of the signal is:

$$y[n] = \sum_{k=1}^p a_k y[n - k] + r[n]$$

# LPC and Speech coding


$$y[n] = \sum_{k=1}^p a_k y[n-k] + r[n]$$

- We only need to send  $a_1, \dots, a_p$  and  $r[n]$ .
- $r[n]$  can be coded with fewer bits



# LPC Example


$$y[n] = \sum_{k=1}^p a_k y[n-k] + r[n]$$



# Infinite Impulse Response (IIR) Systems

In general  $y$  depends on (delayed) samples of the input, as well as the output at previous times (feedback)

$$y[n] = \frac{1}{a_0} \sum_{h=0}^P b_h x[n-h] \quad \leftarrow \text{zeros}$$
$$- \frac{1}{a_0} \sum_{k=1}^Q a_k y[n-k] \quad \leftarrow \text{poles}$$






# Linear Prediction

Auto regressive (AR): only depends on current input and the output at previous times (feedback)


$$y[n] = \frac{b_0}{a_0} x[n] \quad \leftarrow \text{no zeros}$$
$$- \frac{1}{a_0} \sum_{k=1}^Q a_k y[n-k] \quad \leftarrow \text{poles}$$



See also [poles\\_and\\_zeros.pdf](#) in Blackboard

# LPC Limitations

- better match at spectral peaks than at valleys (all-pole model)
- not accurate if transfer function contains zeros (nasals, fricatives...)



# Cepstrum Rationale (Homomorphic Transformation)

- signals combined in a convolutive way:  $a[n] * b[n] * c[n]$
- in the spectral domain:  $A(z)B(z)C(z)$
- taking the log:  $\log(A(z)) + \log(B(z)) + \log(C(z))$
- to analyse the different contribution perform Fourier transform (DCT if not interested in phase information).

# Cepstrum Rationale (Homomorphic Transformation)

- signals combined in a convolutive way:  $a[n] * b[n] * c[n]$
- in the spectral domain:  $A(z)B(z)C(z)$
- taking the log:  $\log(A(z)) + \log(B(z)) + \log(C(z))$
- to analyse the different contribution perform Fourier transform (DCT if not interested in phase information).
- Terminology:
  - frequency vs quefrequency
  - spectrum vs cepstrum
  - filter vs lifter
  - ...

# Cepstrum Definition

## Complex Cepstrum


$$\hat{x}[n] = \frac{1}{2\pi} \int_{-\pi}^{\pi} \ln X(\omega) e^{j\omega n} d\omega$$

## Real Cepstrum

$$c[n] = \frac{1}{2\pi} \int_{-\pi}^{\pi} \ln |X(\omega)| e^{j\omega n} d\omega$$

# Mel Filterbank: Motivation

- Perception of frequencies is logarithmic: Mel and Bark scales
- Perception of amplitude (or energy, or loudness) is logarithmic




$$\text{bark}(f) = 13 \arctan(0.00076 * f) + 3.5 \arctan \left[ \left( \frac{f}{7500} \right)^2 \right]$$


$$\text{mel}(f) = 1125 \ln \left( 1 + \frac{f}{700} \right)$$

# Perceptual Linear Prediction

- Transform to the Bark frequency scale before computing the LPC coefficients
- Cubic root of energy instead of logarithm



# Mel Filterbank: Calculation





# Mel Filterbank: Pros and Cons

## Cons:

- coefficients are correlated (inefficient use of information)
- difficult to model statistically (e.g. multivariate Gaussian distribution)

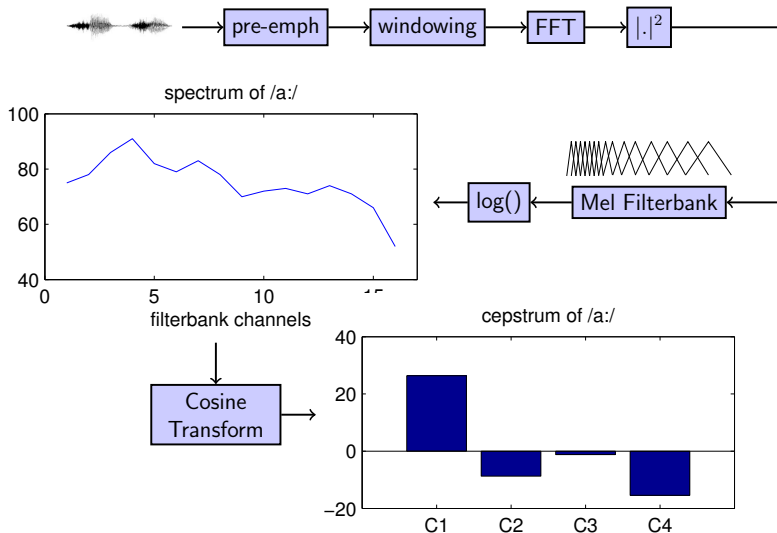
## Pros:

- easy to interpret (smooth spectrum)
- works well with neural networks (no problem with correlations)

# Mel Frequency Cepstrum Coefficients

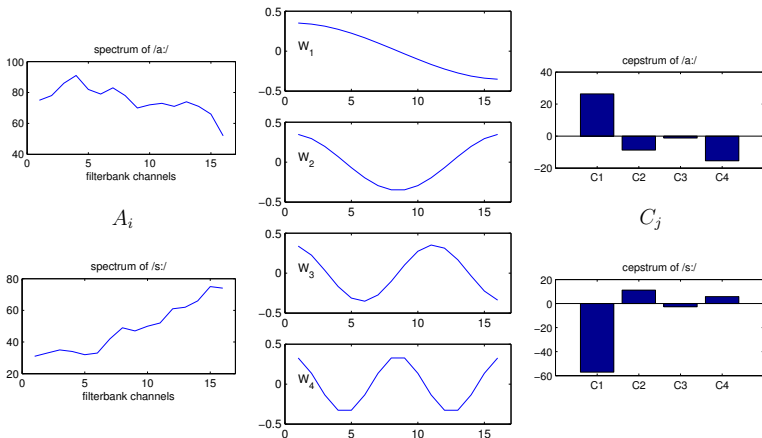
- *de facto* standard in ASR
- imitate aspects of auditory processing
- does not assume all-pole model of the spectrum
- uncorrelated: easier to model statistically

# MFCCs Calculation



# MFCC: Cosine Transform (`scipy.fftpack.realtransforms.dct`)

$$C_j = \sqrt{\frac{2}{N}} \sum_{i=1}^N A_i \cos\left(\frac{j\pi(i-0.5)}{N}\right)$$



# MFCC Advantages

- fairly uncorrelated coefficients (simpler statistical models)
- high phonetic discrimination (empirically shown)
- do not assume all-pole model
- low number of coeff. enough to capture coarse structure of spectrum
- Cepstral Mean Subtraction corresponds to channel removal

# Dynamic Features

Concatenate static MFCCs (or LPCs) to  $\Delta$  and  $\Delta\Delta$  vectors.

$\Delta_n$  computed as weighted sum of  $d_k(n)$

$$\Delta_n = \frac{\sum_{k=1}^K w_k d_k(n)}{\sum_{k=1}^K w_k}$$


$d_k(n)$ : finite differences centered around  $n$  with interval  $2k$ :

$$d_k(n) = \frac{c_{n+k} - c_{n-k}}{2k}$$
$$w_k = 2k^2$$


Similarly for  $\Delta\Delta_n$

# Dynamic Features: Motivation

## Central Differences




## Lanczos Differences



Polynomial fit with or without error

# Dynamic Features: Common values

- Usually  $k$  goes from 1 to 3
- to compute  $\text{static} + \Delta + \Delta\Delta$  we need 13 consecutive static vectors (around 130 msec).





# MFCCs: typical values

- 12 Coefficients C1–C12
- Energy (could be C0)
- Delta coefficients (derivatives in time)
- Delta-delta (second order derivatives)
- total: 39 coefficients per frame (analysis window)